

# Product Requirements Document (PRD): Bible Verse Auto-Detection SaaS Web App (MVP)

## 1. Overview

### 1.1 Product Name

VerseProjection (working title)

### 1.2 Purpose

The VerseProjection SaaS web app continuously analyzes a preacher's or presenter's speech during church services, mapping it to public-domain Bible texts (King James Version, World English Bible) to detect matching verses. When a spoken phrase aligns with a scripture (explicitly, e.g., "John 3:16," or paraphrased, e.g., "God so loved the world"), the app outputs the top 10 best-matching verses in the admin dashboard. Users select the correct verse, which is projected on a screen via a projector. The AI model learns from user selections to improve matching accuracy over time. An optional Electron-based desktop app supports offline use in low-internet environments. The MVP focuses exclusively on Bible verse detection and projection, ensuring simplicity and legal compliance.

### 1.3 Objectives

- Provide real-time, automated projection of Bible verses to enhance sermon delivery.
- Minimize manual effort for church tech teams (primary persona: small/medium church volunteers, ages 25–50).
- Achieve 90%+ detection accuracy and <2-second end-to-end latency (speech to projection).
- Output top 10 matching verses for user selection, supporting both explicit and paraphrased detections.
- Improve model accuracy via user feedback (selections stored for retraining).
- Ensure legal compliance using public-domain KJV/WEB texts.
- Scale for 1,000–10,000 simultaneous users (SOM: 10,000 organizations).

- Launch MVP in 6–9 months within a \$150K–\$400K budget.

## 1.4 Target Audience

- **Primary:** Church worship/tech teams (50–500 members, moderate technical proficiency, urban/suburban).
- **Secondary:** Megachurch media directors (500+ members, high technical expertise, urban).
- **Fringe:** Community/educational religious groups (low budgets, global).

## 1.5 Success Metrics

- **Adoption:** 1,000 organizations in Year 1, 10,000 in 3 years (SOM).
- **Performance:** 90%+ detection accuracy, <2s latency (speech to projection), top 10 matches include correct verse in 95%+ cases.
- **Model Improvement:** 5% accuracy increase after 3 months of user feedback.
- **Revenue:** \$2M/year from 10,000 organizations at \$200/year subscription.
- **User Satisfaction:** 80%+ positive feedback from pilot churches (NPS > 50).
- **Reliability:** 99.9% uptime during peak usage (Sunday services).

# 2. Market and Competitive Context

## 2.1 Market Opportunity

- **TAM:** 1M religious organizations globally using projection software (\$200M at \$200/year).
- **SAM:** 200,000 digitally equipped churches (\$40M).
- **SOM:** 10,000 organizations in 3 years (\$2M).
- **Source:** The Lead Pastor (2023), Barna Group (2020), Capterra (2025).

## 2.2 Competitive Landscape

- **Competitors:** ProPresenter, EasyWorship, MediaShout, OpenLP, Quelea.
- **Differentiator:** Real-time speech analysis for automated Bible verse detection and projection, with top 10 match selection and continuous learning, unlike manual slide-based competitors.
- **Gap:** Competitors lack automation for spontaneous verse references, requiring pre-prepared slides.

## 2.3 Licensing Considerations

- **Bible Texts:** Public-domain KJV and WEB are freely downloadable and storable from repositories like CrossWire Bible Society, Project Gutenberg, or Bible Gateway (JSON exports). No licensing costs or legal risks.
- **Proprietary Bibles:** Deferred to premium tiers (e.g., NIV, ESV via Zondervan, \$1,000–\$10,000/year).
- **Song Lyrics:** Excluded from MVP to avoid CCLI licensing or copyright issues with external APIs or AI models.

## 3. Functional Requirements

### 3.1 Core Features

1. **Real-Time Audio Capture:**
  - Capture continuous audio from a USB microphone or church soundboard via Web Audio API.
  - Support 16-bit, 44.1 kHz audio for high-quality speech detection.
  - Apply client-side noise filtering (e.g., via RecordRTC) to handle church environments (clapping, ambient noise, music).
2. **Verse Detection:**
  - Continuously transcribe preacher's speech using OpenAI Whisper (90%+ accuracy in noisy settings).
  - Map transcribed text to KJV/WEB verses using BERT, identifying matches for explicit references (e.g., "John 3:16") or paraphrased quotes (e.g., "God so loved the world").
  - Output up to 10 best-matching verses ranked by semantic similarity (cosine similarity score, 95%+ chance correct verse in top 10).
  - Process in <1 second (transcription: ~500ms, matching: ~500ms).
  - Support two detection scenarios:
    - **Explicit References:** Preacher cites a verse (e.g., "Please turn to John 3:16"), easily detected via regex and exact matching.
    - **Paraphrased Quotes:** Preacher quotes scripture indirectly (e.g., "God loved the world and gave His Son"), detected via semantic similarity.
3. **Verse Selection and Projection:**
  - Display top 10 matching verses in the admin dashboard with verse text, reference, and confidence score (e.g., "John 3:16 – KJV – 92%").
  - Allow user to click the correct verse for immediate projection on a secondary browser window (full-screen mode) via HDMI/VGA.
  - Support customizable display (font size, color, background).

- Achieve <2-second end-to-end latency (speech to projection, including user selection).
- 4. **Model Learning:**
  - Store user-selected verses in PostgreSQL as feedback data (e.g., transcription, selected verse, confidence scores).
  - Use feedback to incrementally retrain BERT, improving matching accuracy over time (target: 5% accuracy increase after 3 months).
  - Implement periodic retraining (e.g., monthly) via AWS SageMaker.
- 5. **Admin Dashboard:**
  - Configure settings (Bible version: KJV/WEB, display options, audio input).
  - Display real-time transcription and top 10 verse matches with clickable selection.
  - Provide manual override via search bar (e.g., enter "John 3:16") for misdetections.
  - Show detection history and confidence metrics for monitoring.
- 6. **Offline Mode (Desktop App):**
  - Cache KJV/WEB texts (~4MB) and lightweight Whisper/BERT models in Electron app.
  - Perform local audio capture, transcription, and matching (CPU-based).
  - Display top 10 matches for user selection, project verses locally.
  - Store user selections locally, syncing feedback to cloud when online for model retraining.

### 3.2 Non-Functional Requirements

- **Performance:** <2s end-to-end latency, 90%+ detection accuracy, top 10 matches include correct verse in 95%+ cases.
- **Scalability:** Support 1,000–10,000 simultaneous users during peak times (Sunday services).
- **Reliability:** 99.9% uptime via AWS auto-scaling.
- **Security:** HTTPS for communication, JWT for authentication, AES-256 encryption for database and feedback data.
- **Compatibility:** Chrome/Firefox browsers, 720p/1080p projectors, Windows/macOS laptops (8GB RAM, Core i5).
- **Bandwidth:** <500 kbps for audio streaming (Opus codec).
- **Learning:** Model improves by 5% in accuracy after 3 months of user feedback.

## 4. Technical Architecture

## 4.1 Architecture Overview

- **Frontend:** React web app for projection, dashboard, and verse selection; Electron for offline desktop app.
- **Backend:** Node.js/Express for APIs, PostgreSQL for storing KJV/WEB texts and user feedback.
- **AI/ML Pipeline:** OpenAI Whisper for continuous speech-to-text, BERT for semantic verse matching, hosted on AWS SageMaker, optimized with ONNX Runtime.
- **Cloud:** AWS (EC2, RDS, SageMaker, S3, API Gateway, CloudWatch) for hosting, scaling, and logging.
- **Data Flow:** Audio → WebSocket → Backend → AI (transcription, top 10 matches) → Frontend (user selection) → Database (feedback storage) → Frontend (projection).

## 4.2 Connection Flow

1. **Login:**
  - User navigates to [app.verseprojection.com](https://app.verseprojection.com) in Chrome/Firefox on a laptop connected to a projector.
  - Logs in via email/password or SSO (AWS Cognito, JWT authentication).
  - Accesses dashboard to configure settings (e.g., KJV/WEB, font size, audio input).
2. **Service Activation:**
  - User clicks "Start Projection" to initialize WebSocket (Socket.IO) and AI pipeline.
3. **Audio Capture:**
  - Web Audio API captures continuous audio from a USB microphone or soundboard (16-bit, 44.1 kHz).
  - Client-side JavaScript (RecordRTC) applies noise filtering (e.g., high-pass filter at 100 Hz) to reduce background noise (clapping, music).
  - User selects input device in dashboard (dropdown, e.g., "USB Mic", "Soundboard Line-In").
4. **Audio Streaming:**
  - Audio is compressed (Opus codec, <500 kbps, 20ms frames) and streamed via WebSocket to the Node.js backend.
  - AWS API Gateway and Elastic Load Balancer manage concurrent streams from multiple churches.
5. **Audio Analysis:**

- **Transcription:** OpenAI Whisper transcribes audio in real-time (~500ms per 5-second chunk) on AWS SageMaker, fine-tuned for sermon speech (e.g., verse references, religious vocabulary, accents).
- **Verse Matching:**
  - BERT processes transcribed text (~500ms), generating embeddings (e.g., 768-dimensional vectors) for semantic similarity against KJV/WEB verses.
  - For explicit references (e.g., "John 3:16"), apply regex (e.g., `[A-Za-z]+\d+:\d+`) for exact matching, bypassing full semantic analysis (<100ms).
  - For paraphrases, compute cosine similarity between transcription and verse embeddings, ranking top 10 matches (95%+ chance correct verse in list).
  - Output includes verse reference, text, version (KJV/WEB), and confidence score (e.g., 92%).
- Total processing: <1 second (transcription: ~500ms, matching: ~100–500ms).

#### 6. **Verse Selection:**

- Top 10 matches are sent via WebSocket to the React frontend and displayed in the admin dashboard (e.g., table with columns: Reference, Text, Version, Confidence).
- User clicks the correct verse (e.g., "John 3:16 – KJV"), triggering immediate projection.
- Selection is logged in PostgreSQL (transcription, selected verse, confidence scores) for model retraining.

#### 7. **Projection:**

- Selected verse is rendered on a secondary browser window (full-screen, 720p/1080p) connected to a projector via HDMI/VGA.
- Total latency (speech to projection, including user selection): <2 seconds.

#### 8. **Monitoring and Overrides:**

- Dashboard shows real-time transcription, top 10 matches, and selection history.
- Manual override via search bar (e.g., type "John 3:16") for misdetections.
- AWS CloudWatch logs errors (e.g., low-confidence matches, transcription failures) for debugging.

#### 9. **Model Learning:**

- User selections are stored in PostgreSQL (**feedback** table).

- Monthly retraining job on AWS SageMaker updates BERT weights using feedback data (e.g., transcription → selected verse pairs), improving similarity scoring.

## 4.3 Key Modules

### 1. Audio Input Module:

- **Technologies:** Web Audio API, PyAudio (desktop app), RecordRTC, OpenAI Whisper.
- **Function:** Captures and preprocesses continuous audio from microphone or soundboard.
- **Requirements:**
  - 16-bit/44.1 kHz quality for clear speech.
  - Noise filtering (e.g., 100 Hz high-pass, 8 kHz low-pass) for church environments.
  - <200ms capture latency, <500 kbps streaming bandwidth.

### 2. NLP/Text Matching Module:

- **Technologies:** BERT (Hugging Face Transformers), PostgreSQL, ONNX Runtime, regex for explicit references.
- **Function:**
  - Transcribes speech via Whisper (90%+ accuracy).
  - Matches text to KJV/WEB verses, outputting top 10 matches (95%+ correct verse inclusion).
  - Handles explicit references (regex, <100ms) and paraphrases (semantic similarity, ~500ms).
- **Requirements:**
  - <1-second processing (transcription: ~500ms, matching: ~100–500ms).
  - Fine-tuned for sermon speech and verse paraphrases.
  - Confidence scores for ranking (e.g., cosine similarity 0–1).

### 3. Database Module:

- **Technologies:** PostgreSQL with full-text search and vector extensions (e.g., pgvector for embeddings).
- **Function:** Stores KJV/WEB verses, user settings, and feedback data for retraining.
- **Requirements:**
  - ~4MB for KJV/WEB texts, sub-100ms query latency.
  - Supports vector similarity searches for BERT embeddings.
  - Sourced from public-domain repositories (e.g., CrossWire Bible Society).

- AES-256 encryption for data at rest.
- 4. **Projection Display Module:**
  - **Technologies:** React, Socket.IO, Electron (offline).
  - **Function:** Renders user-selected verse on projector (720p/1080p) via full-screen browser window.
  - **Requirements:**
    - Customizable (font size: 12–48pt, color, background).
    - <200ms rendering latency.
    - Compatible with HDMI/VGA projectors.
- 5. **Admin Dashboard Module:**
  - **Technologies:** React, Material-UI.
  - **Function:**
    - Configures settings (Bible version, display, audio input).
    - Displays real-time transcription and top 10 verse matches (clickable table).
    - Provides manual override and detection history.
  - **Requirements:**
    - Intuitive for non-technical users (drag-and-drop settings, clear buttons).
    - Real-time updates via WebSocket.
- 6. **Cloud Integration Module:**
  - **Technologies:** AWS (EC2, RDS, SageMaker, S3, API Gateway, CloudWatch).
  - **Function:** Hosts backend, AI, database; scales for peak usage; manages retraining.
  - **Requirements:**
    - 99.9% uptime.
    - \$400–\$800/month cost for 1,000 users.
    - Auto-scaling for 1,000–10,000 users.

## 5. Skillset Requirements

### 5.1 Roles and Responsibilities

1. **AI/ML Engineer:**
  - **Skills:** Python, TensorFlow/PyTorch, OpenAI Whisper, BERT, ONNX, audio processing (Librosa), incremental learning.
  - **Tasks:**
    - Develop/fine-tune Whisper for sermon transcription.



- Fine-tune BERT for verse matching (explicit and paraphrased).
  - Implement top 10 match ranking and feedback-based retraining.
  - Optimize pipeline for <1s latency using ONNX.
- **Need:** Critical for detection accuracy and model learning.
- 2. **Backend Developer:**
  - **Skills:** Node.js/Express, PostgreSQL (pgvector), AWS (EC2, RDS, API Gateway).
  - **Tasks:**
    - Build APIs for audio streaming, verse matching, and feedback storage.
    - Manage database (KJV/WEB, feedback data).
    - Ensure scalability and security (JWT, AES-256).
  - **Need:** Essential for system reliability and data handling.
- 3. **Frontend Developer:**
  - **Skills:** React, Socket.IO, Material-UI/Tailwind CSS, Electron, Web Audio API.
  - **Tasks:**
    - Develop dashboard with top 10 match selection and override functionality.
    - Build projection UI and offline Electron app.
    - Ensure responsive, user-friendly design.
  - **Need:** Critical for user experience and adoption.
- 4. **DevOps Engineer (Part-Time):**
  - **Skills:** AWS (EC2, RDS, SageMaker, CloudWatch), CI/CD (GitHub Actions).
  - **Tasks:**
    - Deploy app to AWS, optimize costs (\$400–\$800/month).
    - Set up auto-scaling and monitoring.
    - Manage SageMaker retraining jobs.
  - **Need:** Necessary for scalability and operational efficiency.
- 5. **Product Manager:**
  - **Skills:** Church technology knowledge, MVP scoping, user feedback analysis.
  - **Tasks:**
    - Define features (multi-match, learning).
    - Source public-domain Bible texts (e.g., CrossWire).
    - Coordinate pilot testing with churches.
  - **Need:** Essential for market alignment and user-centric design.

## 5.2 Team Size

- **Core Team:** 3–4 members (AI/ML Engineer, Full-Stack Developer, Part-Time DevOps, Product Manager).
- **Timeline:** 6–9 months (slightly increased complexity for multi-match and learning).
- **Budget:** \$150K–\$400K (salaries, AWS, tools).

## 6. Implementation Details

### 6.1 Technology Stack

- **Frontend:**
  - React (UI framework for dashboard, projection, and match selection).
  - Socket.IO (real-time WebSocket for audio, matches, and projection).
  - Material-UI/Tailwind CSS (responsive, user-friendly styling).
  - Electron (offline desktop app).
- **Backend:**
  - Node.js/Express (API server for audio, matching, feedback).
  - PostgreSQL (database for KJV/WEB, feedback, users).
- **AI/ML:**
  - OpenAI Whisper (continuous speech-to-text, fine-tuned for sermons).
  - BERT (Hugging Face Transformers, semantic matching, top 10 ranking).
  - TensorFlow/PyTorch (model training frameworks).
  - Librosa (audio preprocessing, e.g., noise reduction).
  - ONNX Runtime (low-latency inference).
- **Cloud:**
  - AWS EC2 (backend hosting, t3.medium, auto-scaling).
  - AWS RDS (PostgreSQL, db.t3.micro, 5GB storage).
  - AWS SageMaker (AI inference and retraining, ml.t3.medium).
  - AWS S3 (logs, feedback data backups).
  - AWS API Gateway (WebSocket management).
  - AWS CloudWatch (error logging, performance monitoring).
- **Tools:**
  - GitHub (version control).
  - GitHub Actions (CI/CD pipeline).
  - RecordRTC (client-side audio preprocessing).
  - Web Audio API (browser-based audio capture).

### 6.2 Database Schema

- **Bible Table:**

- **Columns:**
  - **id** (UUID, primary key).
  - **version** (string, e.g., "KJV", "WEB").
  - **book** (string, e.g., "John").
  - **chapter** (integer, e.g., 3).
  - **verse** (integer, e.g., 16).
  - **text** (text, e.g., "For God so loved the world...").
  - **embedding** (vector, 768-dimensional BERT embedding for similarity).
- **Size:** ~4MB for KJV/WEB texts, ~8MB with embeddings.
- **Indexing:**
  - Full-text search index on **text** for manual overrides.
  - Vector index (pgvector) on **embedding** for semantic similarity (cosine distance).
- **Source:** Public-domain JSON exports from CrossWire Bible Society or Project Gutenberg.
- **Security:** AES-256 encryption.
- **Users Table:**
  - **Columns:**
    - **id** (UUID, primary key).
    - **email** (string, unique).
    - **password\_hash** (string, bcrypt).
    - **subscription\_tier** (string, e.g., "free", "premium").
    - **settings** (JSON, e.g., [{"bible\_version": "KJV", "font\_size": 24}]).
  - **Security:** AES-256 encryption, JWT for sessions.
- **Feedback Table:**
  - **Columns:**
    - **id** (UUID, primary key).
    - **user\_id** (UUID, foreign key to Users).
    - **timestamp** (datetime).
    - **transcription** (text, e.g., "God loved the world").
    - **selected\_verse\_id** (UUID, foreign key to Bible).
    - **top\_matches** (JSON, e.g., [{"verse\_id": "uuid", "confidence": 0.92}, ...]).
  - **Size:** ~10MB/month for 1,000 users (assuming 100 selections/user/month).
  - **Use:** Stores data for BERT retraining, improving accuracy.

- **Security:** AES-256 encryption.
- **Logs Table** (optional, for debugging):
  - **Columns:**
    - **id** (UUID, primary key).
    - **timestamp** (datetime).
    - **user\_id** (UUID, foreign key).
    - **event** (string, e.g., "low\_confidence\_match").
    - **details** (JSON, e.g., {"transcription": "text", "confidence": 0.6}).
  - **Size:** ~1MB/month for 1,000 users.
  - **Storage:** AWS S3 or CloudWatch.

### 6.3 AI Pipeline

- **Input:** Continuous audio stream (16-bit, 44.1 kHz, Opus codec, 5-second chunks).
- **Processing:**
  - **Whisper:**
    - Transcribes audio to text (~500ms per chunk).
    - Fine-tuned on sermon audio datasets (~10 hours, public-domain or synthetic) for religious vocabulary, verse references, and noise resilience (clapping, music).
    - Parameters: Whisper-small (~240M parameters, ~500MB memory).
  - **BERT:**
    - Processes transcribed text (~500ms), generating 768-dimensional embeddings.
    - For explicit references (e.g., "John 3:16"):
      - Regex pattern `[A-Za-z]+\d+:\d+` extracts verse references.
      - Exact match against Bible table (<100ms).
    - For paraphrases:
      - Compute cosine similarity between transcription embedding and precomputed verse embeddings (stored in **Bible.embedding**).
      - Rank top 10 matches by similarity score (0–1, e.g., 0.92 for "John 3:16").
      - Cache embeddings in memory (Redis, ~8MB) for sub-100ms queries.

- Parameters: BERT-base-uncased (~110M parameters, ~400MB memory).
- **Optimization:**
  - ONNX Runtime reduces inference time (e.g., 30% faster on CPU).
  - Batch processing for overlapping audio chunks (e.g., 5s chunks every 2s).
- **Output:**
  - Array of top 10 matches: `[[verse_id, reference, text, version, confidence], ...]`.
  - Sent to frontend via WebSocket.
- **Feedback Loop:**
  - User selections stored in **Feedback** table (transcription, selected verse, top matches).
  - Monthly SageMaker job retrains BERT using feedback data (~1,000–10,000 pairs/month for 1,000 users).
  - Loss function: Triplet loss to optimize embedding distances (correct verse closer, incorrect farther).
  - Target: 5% accuracy increase after 3 months.
- **Training Data:**
  - Whisper: ~10 hours of sermon audio (public-domain sermons, synthetic clips with verse references).
  - BERT: ~10,000 transcription-verse pairs (KJV/WEB texts, paraphrased examples, e.g., "God so loved" → John 3:16).
  - Feedback: Incremental user selections (e.g., 1,000 pairs/week for 1,000 users).
- **Hosting:** AWS SageMaker (ml.t3.medium instances, auto-scaling).

## 6.4 Cloud Infrastructure

- **AWS EC2:**
  - Hosts Node.js/Express backend.
  - Instance: t3.medium (2 vCPUs, 4GB RAM).
  - Auto-scaling: 1–10 instances for Sunday peaks.
  - Cost: ~\$50–\$200/month for 1,000 users.
- **AWS RDS:**
  - PostgreSQL database (KJV/WEB, feedback, users).
  - Instance: db.t3.micro (1 vCPU, 1GB RAM, 10GB storage for ~12MB data).
  - Extensions: pgvector for vector similarity searches.
  - Cost: ~\$50–\$100/month.
- **AWS SageMaker:**

- Hosts Whisper and BERT for inference and retraining.
- Instance: ml.t3.medium (2 vCPUs, 4GB RAM) for inference.
- Retraining: ml.m5.large (2 vCPUs, 8GB RAM, ~1 hour/month).
- Cost: ~\$250–\$450/month (inference: \$200–\$400, retraining: \$50).
- **AWS S3:**
  - Stores logs, feedback backups, and optional audio recordings.
  - Cost: ~\$10–\$50/month.
- **AWS API Gateway:**
  - Manages WebSocket connections (audio, matches, selections).
  - Cost: ~\$20–\$50/month.
- **AWS CloudWatch:**
  - Logs errors (e.g., low-confidence matches, transcription failures).
  - Monitors latency, accuracy, and user selection rates.
  - Cost: ~\$10–\$20/month.
- **AWS ElastiCache (Redis):**
  - Caches verse embeddings (~8MB) for sub-100ms similarity queries.
  - Instance: cache.t3.micro (1 vCPU, 0.5GB RAM).
  - Cost: ~\$20–\$30/month.
- **Total Cost:** \$460–\$900/month for 1,000 users, optimized with Reserved Instances.

## 6.5 Offline Mode (Desktop App)

- **Technologies:** Electron, lightweight Whisper/BERT models, SQLite, PyAudio.
- **Storage:**
  - ~4MB for KJV/WEB texts, ~8MB for precomputed embeddings (SQLite).
  - ~500MB for Whisper-small and BERT-base models (ONNX format).
- **Processing:**
  - Audio capture via PyAudio (16-bit, 44.1 kHz).
  - Local Whisper/BERT inference (CPU-based, 8GB RAM, Core i5).
  - Top 10 match ranking using SQLite vector searches (pgvector equivalent).
- **Functionality:**
  - Displays top 10 matches in Electron UI for user selection.
  - Projects selected verse in full-screen window.
  - Stores selections locally (SQLite **feedback** table).
  - Syncs feedback to cloud (PostgreSQL) when online for retraining.
- **Requirements:**
  - ~600MB disk space.
  - <2-second latency (local processing).

- No internet dependency.

## 7. User Experience

### 7.1 User Flow

1. **Setup:**
  - Connect laptop to projector (HDMI/VGA).
  - Log into web app ([app.verseprojection.com](https://app.verseprojection.com)) or launch Electron app (offline).
2. **Configuration:**
  - Select KJV or WEB, audio input (mic/soundboard), display settings (font size: 12–48pt, color, background).
  - Adjust sensitivity (e.g., confidence threshold: 0.7–0.9) for match detection.
3. **Operation:**
  - Click “Start Projection”; app continuously analyzes preacher’s speech.
  - Top 10 matching verses appear in dashboard (e.g., “John 3:16 – KJV – 92%”).
  - User clicks correct verse for projection (e.g., ~1-second decision).
  - Verse displays on projector in full-screen window.
4. **Overrides:**
  - If no correct match, use search bar to manually select verse (e.g., “John 3:16”).
  - Dashboard shows transcription and detection history for context.
5. **Offline:**
  - Electron app runs locally, mirroring web app functionality.
  - Syncs selections to cloud when reconnected.

### 7.2 UI Components

- **Login Page:**
  - Email/password or SSO (AWS Cognito).
  - Minimal design, dark/light theme toggle.
- **Dashboard:**
  - **Settings Panel:**
    - Dropdowns: Bible version (KJV/WEB), audio input, projector resolution (720p/1080p).
    - Sliders: Font size (12–48pt), confidence threshold (0.7–0.9).
    - Color picker: Text, background.

- **Transcription View:**
  - Real-time text (e.g., "Pastor: 'God so loved the world'").
  - Scrollable, 10-second history.
- **Matches Table:**
  - Columns: Reference (e.g., "John 3:16"), Text (preview, e.g., "For God so loved..."), Version (KJV/WEB), Confidence (e.g., 92%).
  - Clickable rows to select verse for projection.
  - Sorted by confidence (descending).
- **Override Search Bar:**
  - Autocomplete for verse references (e.g., type "Jo" → "John").
  - Instant display on selection.
- **Detection History:**
  - Timeline of past matches and selections (e.g., "10:02 AM: John 3:16 selected").
- **Projection Window:**
  - Full-screen browser or Electron window.
  - Customizable verse text (e.g., white text, black background, 24pt Arial).
  - Smooth transitions (200ms fade-in).
- **Error Notifications:**
  - Pop-ups for issues (e.g., "Low audio quality, check mic").
  - Suggests manual override for low-confidence matches (<0.7).

## 7.3 UX Principles

- **Simplicity:** One-click projection start, intuitive match selection (clickable table).
- **Reliability:** Real-time feedback (transcription, matches), robust overrides.
- **Accessibility:**
  - High-contrast text (WCAG 2.1 compliant).
  - Keyboard navigation (e.g., arrow keys for match selection).
  - Support for 720p/1080p projectors.

## 8. Constraints and Risks

### 8.1 Technical Constraints

- **Latency:** Must achieve <2s end-to-end latency (speech to projection, including user selection).
- **Noise:** Church environments (clapping, music, congregation) may reduce Whisper's transcription accuracy.



- **Multi-Match Accuracy:** Top 10 matches must include correct verse in 95%+ cases, requiring robust BERT fine-tuning.
- **Browser Compatibility:** Web Audio API and WebSocket vary across browsers (e.g., Safari limitations).
- **Hardware:** Laptops need 8GB RAM, Core i5; projectors require 720p/1080p support.
- **Learning:** Feedback loop must improve accuracy (5% in 3 months) without degrading performance.

## 8.2 Licensing Risks

- **Bibles:** Public-domain KJV/WEB texts are legally free to download and store (e.g., CrossWire Bible Society). No risks.
- **Proprietary Bibles:** Deferred to premium tiers to avoid licensing costs (\$1,000–\$10,000/year).

## 8.3 Mitigation Strategies

- **Noise:**
  - Fine-tune Whisper on sermon audio with noise (clapping, music, ~10 hours).
  - Apply advanced filtering (Librosa, high-pass at 100 Hz, low-pass at 8 kHz).
  - Provide manual override for transcription errors.
- **Multi-Match Accuracy:**
  - Fine-tune BERT on ~10,000 transcription-verse pairs, emphasizing paraphrases.
  - Use feedback data to retrain monthly, prioritizing high-confidence mismatches.
  - Set default confidence threshold (0.7) to balance sensitivity and precision.
- **Browser:**
  - Optimize for Chrome (80% church usage, Capterra 2025), test Firefox.
  - Provide setup guides for audio/WebSocket troubleshooting.
  - Fallback to Electron app for incompatible browsers.
- **Hardware:**
  - Support multiple projector resolutions (720p, 1080p).
  - Recommend HDMI/VGA adapters in user documentation.
  - Test Electron app on low-end laptops (8GB RAM, Core i5).
- **Learning:**

- Validate feedback data (e.g., filter low-confidence selections).
- Monitor retraining impact via CloudWatch (accuracy, latency).
- Cap retraining to ~1 hour/month to control costs.
- **Licensing:**
  - Source KJV/WEB from verified repositories (CrossWire Bible Society, Project Gutenberg).
  - Document sourcing in app (e.g., "KJV sourced from CrossWire, public domain").
- **Tarpit Risk:**
  - Focused scope (verses only) and public-domain texts minimize complexity.
  - Continuous analysis and multi-match selection add moderate complexity, mitigated by leveraging existing AI pipeline and robust testing.

## 9. Development Timeline

- **Month 1–2:**
  - Gather requirements, source KJV/WEB texts (CrossWire Bible Society).
  - Set up PostgreSQL database (Bible, Feedback tables).
  - Collect sermon audio for Whisper fine-tuning (~10 hours).
- **Month 3–4:**
  - Develop AI pipeline:
    - Fine-tune Whisper for sermon speech (noise, accents).
    - Fine-tune BERT for verse matching (explicit references, paraphrases).
    - Implement top 10 match ranking and feedback storage.
  - Build backend APIs (Node.js/Express) for audio, matching, feedback.
- **Month 5–6:**
  - Develop frontend (React):
    - Dashboard with match selection table, override search.
    - Projection window with customizable display.
  - Integrate WebSocket (Socket.IO) for real-time updates.
- **Month 7–8:**
  - Deploy to AWS (EC2, RDS, SageMaker, API Gateway, CloudWatch, ElastiCache).
  - Develop offline Electron app (SQLite, local Whisper/BERT).
  - Test in simulated church environments (noise, paraphrases, explicit references).

- **Month 8–9:**
  - Pilot with 5–10 churches (urban/suburban, 50–500 members).
  - Collect feedback data (~1,000 selections/pilot).
  - Iterate UI (e.g., simplify match selection) and model (first retraining).
  - Finalize MVP for launch.

## 10. Budget Allocation

- **Salaries:** \$120K–\$300K (3–4 team members, 6–9 months).
  - AI/ML Engineer: \$50K–\$100K.
  - Full-Stack Developer: \$50K–\$100K.
  - Part-Time DevOps: \$10K–\$30K.
  - Product Manager: \$10K–\$30K.
- **AWS Hosting:** \$6K–\$18K (EC2, RDS, SageMaker, S3, API Gateway, CloudWatch, ElastiCache).
  - Monthly: \$460–\$900 for 1,000 users.
  - Annual: ~\$5.5K–\$10.8K.
- **Tools/Software:** \$5K (GitHub, Hugging Face, testing tools, audio datasets).
- **Testing/Pilot:** \$2K–\$5K (church outreach, feedback collection).
- **Total:** \$153K–\$428K (slightly increased due to feedback loop and multi-match complexity, within \$150K–\$400K target via cost optimization).

## 11. Sources

- Capterra (2025): Church software trends.
- Barna Group (2020): Church technology preferences.
- AWS Documentation (2023): Cloud and AI best practices.
- Hugging Face (2023): Whisper and BERT model specifications.
- MDN (2023): Web Audio API and WebSocket documentation.
- CrossWire Bible Society (2023): Public-domain Bible texts (KJV, WEB).
- U.S. Copyright Office (2023): Public domain guidelines.
- Librosa Documentation (2023): Audio preprocessing techniques.
- ONNX Runtime (2023): Model optimization for inference.

## 12. Conclusion

The VerseProjection MVP delivers real-time Bible verse auto-detection and projection, continuously analyzing preacher speech to map explicit or paraphrased quotes to public-domain KJV/WEB texts. The app outputs the top 10 matching

verses for user selection, projects the chosen verse, and learns from selections to improve accuracy (5% increase in 3 months). The technical architecture (Web Audio API → Whisper → BERT → PostgreSQL → React, hosted on AWS) ensures <2-second latency, 90%+ accuracy, and scalability for 1,000–10,000 users. An offline Electron app supports low-internet environments. By focusing on verses only, the MVP avoids licensing complexities, staying within a \$150K–\$400K budget and 6–9-month timeline. The solution targets 10,000 churches (\$2M revenue), addressing the need for automated, user-friendly verse projection during sermons.